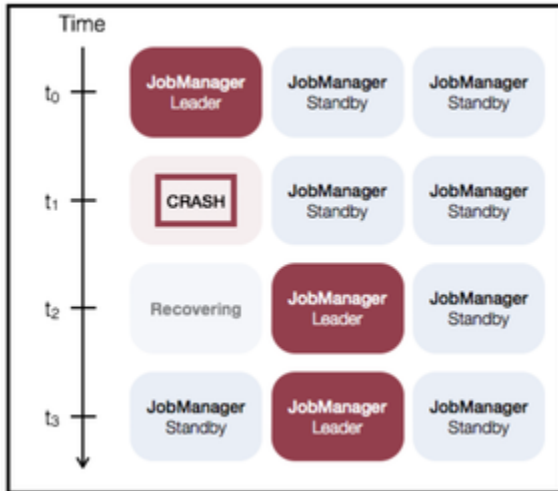


Install Apache Flink on Multi-node Cluster

Architecture

The general idea of JobManager high availability for standalone clusters is that there is a **single leading JobManager** at any time and **multiple standby JobManagers** to take over leadership in case the leader fails. This guarantees that there is **no single point of failure** and programs can make progress as soon as a standby JobManager has taken leadership. There is no explicit distinction between standby and master JobManager instances. Each JobManager can take the role of master or standby.

As an example, consider the following setup with three JobManager instances:



Platform Requirements

1. Operating system: Ubuntu 16.04 or later.
2. Java 8.x or higher.
3. Create an odd number of machines (eg: 3, 5, 7..).
4. System requirement is job specific.
 - a. Minimum memory requirement is 4 GB/ machine.
 - b. Select the number of CPU cores according to the number of flink jobs.
5. Need a shared filesystem. For example.
 - a. Elastic File System (EFS, Managed service provided by AWS.)
 - b. Gluster File System (GFS)
 - c. Hadoop Distributed File System (HDFS)

Prerequisites

The below document creates a 3 node cluster

1. Install Java 8 (For each Flink machines).

```
apt-get update && apt-get -y upgrade
apt-get install openjdk-8-jdk-headless
```

2. Setup environment (For each Flink machines).

```
echo ""
LANGUAGE=\"en_US.UTF-8\"
LANG=\"en_US.UTF-8\"
LC_ALL=\"en_US.UTF-8\"
JAVA_HOME=\"/usr/lib/jvm/java-8-openjdk-amd64\"
"" >> /etc/environment
```

3. Assign proper hostname for machines (For each Flink machines).

```
eg: flink-server1
hostnamectl set-hostname flink-server1
hostname -F /etc/hostname
echo ""
<Ip of flink-server1> flink-server1
<Ip of flink-server2> flink-server2
<Ip of flink-server3> flink-server3
"" > /etc/hosts
```

4. Configure passwordless root user ssh between each machine.

```
eg: flink-server1

Generate the ssh-keypair
sudo su -
apt-get install openssh-server openssh-client
ssh-keygen -t rsa -P ""
cat /root/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys

Copy the public key From "flink-server1" (cat
/root/.ssh/id_rsa.pub) and add that key in other servers
/root/.ssh/authorized_keys.

Test passwordless ssh, try root ssh from flink-server1 to
flink-server2 & flink-server3 and vice-versa
```

Installation and configuration

1. Installation => Take 1 machine (eg: flink-server1)

```
cd /opt/  
wget  
https://archive.apache.org/dist/flink/flink-1.7.0/flink-1.7.0-bin-scala_2.12.tgz  
tar -zxvf flink-1.7.0-bin-scala_2.12.tgz  
cd flink-1.7.0
```

2. Configuration => Take 1 machine (eg: flink-server1)
 - a. Edit /opt/flink-1.7.0/conf/flink-conf.yaml (Main configuration file)

```
Comment the below line  
# jobmanager.rpc.address: localhost  
  
Assign necessary heap memory as required  
jobmanager.heap.size: 3072m  
taskmanager.heap.size: 8192m  
  
Assign necessary task slots as required (recommended assign  
number_of_cpu_cores-1 slots)  
taskmanager.numberOfTaskSlots: 7  
  
Add the below entries at the end of this configuration file  
high-availability: zookeeper  
high-availability.zookeeper.quorum:  
flink-server1:2181,flink-server1:2181,flink-server3:2181  
high-availability.cluster-id: /cluster_one  
high-availability.storageDir: file:///opt/flink-1.7.0/ha  
zookeeper.sasl.disable: true  
  
akka.watch.heartbeat.pause: 600s  
akka.watch.threshold: 50  
  
akka.ask.timeout: 600s  
  
restart-strategy: failure-rate  
restart-strategy.failure-rate.max-failures-per-interval: 10  
restart-strategy.failure-rate.failure-rate-interval: 5 min  
restart-strategy.failure-rate.delay: 10 s
```

- b. Edit /opt/flink-1.7.0/conf/zoo.cfg (Zookeeper configuration file)

Add below ZooKeeper quorum peers by

Removing the entry

```
server.1=localhost:2888:3888
```

Adding the entry

```
server.1=flink-server1:2888:3888
```

```
server.2=flink-server2:2888:3888
```

```
server.3=flink-server3:2888:3888
```

c. Edit /opt/flink-1.7.0/conf/masters (Flink master configuration file)

Add the below entries

```
flink-server1:8081
```

```
flink-server2:8081
```

```
flink-server3:8081
```

d. Edit /opt/flink-1.7.0/conf/slaves (Flink slave configuration file)

Add the below entries

```
flink-server1
```

```
flink-server2
```

```
flink-server3
```

e. Edit /opt/flink-1.7.0/conf/log4j.properties (Flink log configuration file)

```
Comment the below lines
#log4j.appender.file=org.apache.log4j.FileAppender
#log4j.appender.file.file=${log.file}
#log4j.appender.file.append=false
#log4j.appender.file.layout=org.apache.log4j.PatternLayout
#log4j.appender.file.layout.ConversionPattern=%d{yyyy-MM-dd
HH:mm:ss,SSS} %-5p %-60c %x - %m%n
```

```
Adding the below lines at the end of configuration file.
log4j.appender.file=org.apache.log4j.RollingFileAppender
log4j.appender.file.file=${log.file}
log4j.appender.file.MaxFileSize=300MB
log4j.appender.file.MaxBackupIndex=10
log4j.appender.file.append=false
log4j.appender.file.layout=org.apache.log4j.PatternLayout
log4j.appender.file.layout.ConversionPattern=%d{yyyy-MM-dd
HH:mm:ss,SSS} %-5p %-60c %x - %m%n
```

f. Copy flink installation folder to other machines

```
scp -r /opt/flink-1.7.0 flink-server2:/opt
scp -r /opt/flink-1.7.0 flink-server3:/opt
```

3. Create and configure the shared file system

- a. Create EFS [Link](#), [more details](#).
- b. Mount the filesystem in each machine.

```
Mounting, execute the below commands.
mkdir /opt/flink-1.7.0/ha
echo "<file_system_id>.efs.us-east-1.amazonaws.com:/
/opt/flink-1.7.0/ha nfs defaults,_netdev 0 0" >> /etc/fstab
mount -a
```

```
Confirm the mount, execute the below command.
df -Th
```

Flink cluster managing commands

1. Starting the cluster

```
Login to any flink server and change to flink installation folder.  
cd /opt/flink-1.7.0/
```

```
Start zookeeper quorum  
./bin/start-zookeeper-quorum.sh
```

```
Starting the flink cluster  
./bin/start-cluster.sh
```

2. Stopping the cluster

```
Login to any flink server and change to flink installation folder.  
cd /opt/flink-1.7.0/
```

```
Start zookeeper quorum  
./bin/stop-cluster.sh
```

```
Starting the flink cluster  
./bin/stop-zookeeper-quorum.sh
```

3. Starting JobManager and TaskManager

```
Login to any flink server and change to flink installation folder.  
cd /opt/flink-1.7.0/
```

```
Starting job manager  
./bin/jobmanager.sh start
```

```
Starting task manager  
./bin/taskmanager.sh start
```

4. Stopping JobManager and TaskManager

```
Login to any flink server and change to the flink installation  
folder.  
cd /opt/flink-1.7.0/
```

```
Starting job manager  
./bin/jobmanager.sh stop
```

```
Starting task manager  
./bin/taskmanager.sh stop
```

Note: AutoScaling for both jobmanager and taskmanager is possible by adding a new server

- Copy the flink folder from one of the servers to the new server in the same path.
- EFS mount is not required for taskmanager autoscaling
- Run command to start jobmanager/ taskmanager.

Accessing Flink Servers

Flink master provides a UI listening on **TCP port 8081** to manage the resources through a REST interface. Since we have multiple masters, it is required to create a LoadBalancers with proper health check for accessing Flink UI.

In the case of AWS,

1. It is recommended to use an internal Classic Load Balancer for internal routing.
2. For external access use either VPN (Secure access for prod and stage) or a public Classic Load Balancer.

Flink AWS CLB health check configuration is shown below.

Ping Protocol	HTTP
Ping Port	8081
Ping Path	/
Advanced Details	
Response Timeout	3 seconds
Interval	5 seconds
Unhealthy threshold	2
Healthy threshold	2

Monitoring and Recovery

1. Cronjob monitoring and Recovery.

Assign the cronjob for each machine.

Commands

```
cd /opt/flink-1.7.0/
```

```
mkdir custom
```

```
vim /opt/flink-1.7.0/custom/autorestart.sh
```

Add the code =>

```
https://github.com/xblockchainlabs/csfx-cicd/blob/master/sources/Flink-Job-Monitoring/flink-service-restart-cron.sh
```

```
chmod +x /opt/flink-1.7.0/custom/autorestart.sh
```

```
crontab -e
```

Add the below lines

```
* * * * * bash /opt/flink-1.7.0/custom/autorestart.sh
```

```
* * * * * ( sleep 30 ; bash /opt/flink-1.7.0/custom/autorestart.sh)
```

Verify cronjobs

```
crontab -l
```

2. AWS lambda monitoring for flink jobs.

```
a) The lambda function is using SES for sending emails, verify the
sender and receiver email addresses before creating the Lambda
function
b) Create a lambda function with python3.7 run time and assign
proper IAM roles to access NIC and cloudwatch.
c) Assign proper subnets in the vpc so that the lambda can access
flink servers
d) Add the below environment variables.
ENVIRONMENT => eg: Test/ DEV
EXPECTEDJOBS => eg:
ohlcv24h,ohlcv8h,ohlcv30m,rollup1m,ohlcv1m,ohlcv5m,ohlcv15m,rollupm
ins,ohlcv72h,ohlcv1h,rolluphrs,ohlcv48h
FLINKSERVER => Internal LoadBalancer URL
RECIPIENTS => Recipients email addresses
SENDER => senders email address
e) Use python virtual env to package the below script. Zip the code
block into one file and upload it to lambda function.
f) Use the below handler for the lambda function.
flinkJobMonitoring.flinkJobMonitoring
```

3. Install Logstash and push logs into ElasticSearch for enabling centralized log monitoring.

```
Execute the below bash script in each servers.
https://github.com/xblockchainlabs/csfx-cicd/blob/master/sources/Flink-Job-Monitoring/flink-logstash.sh
```